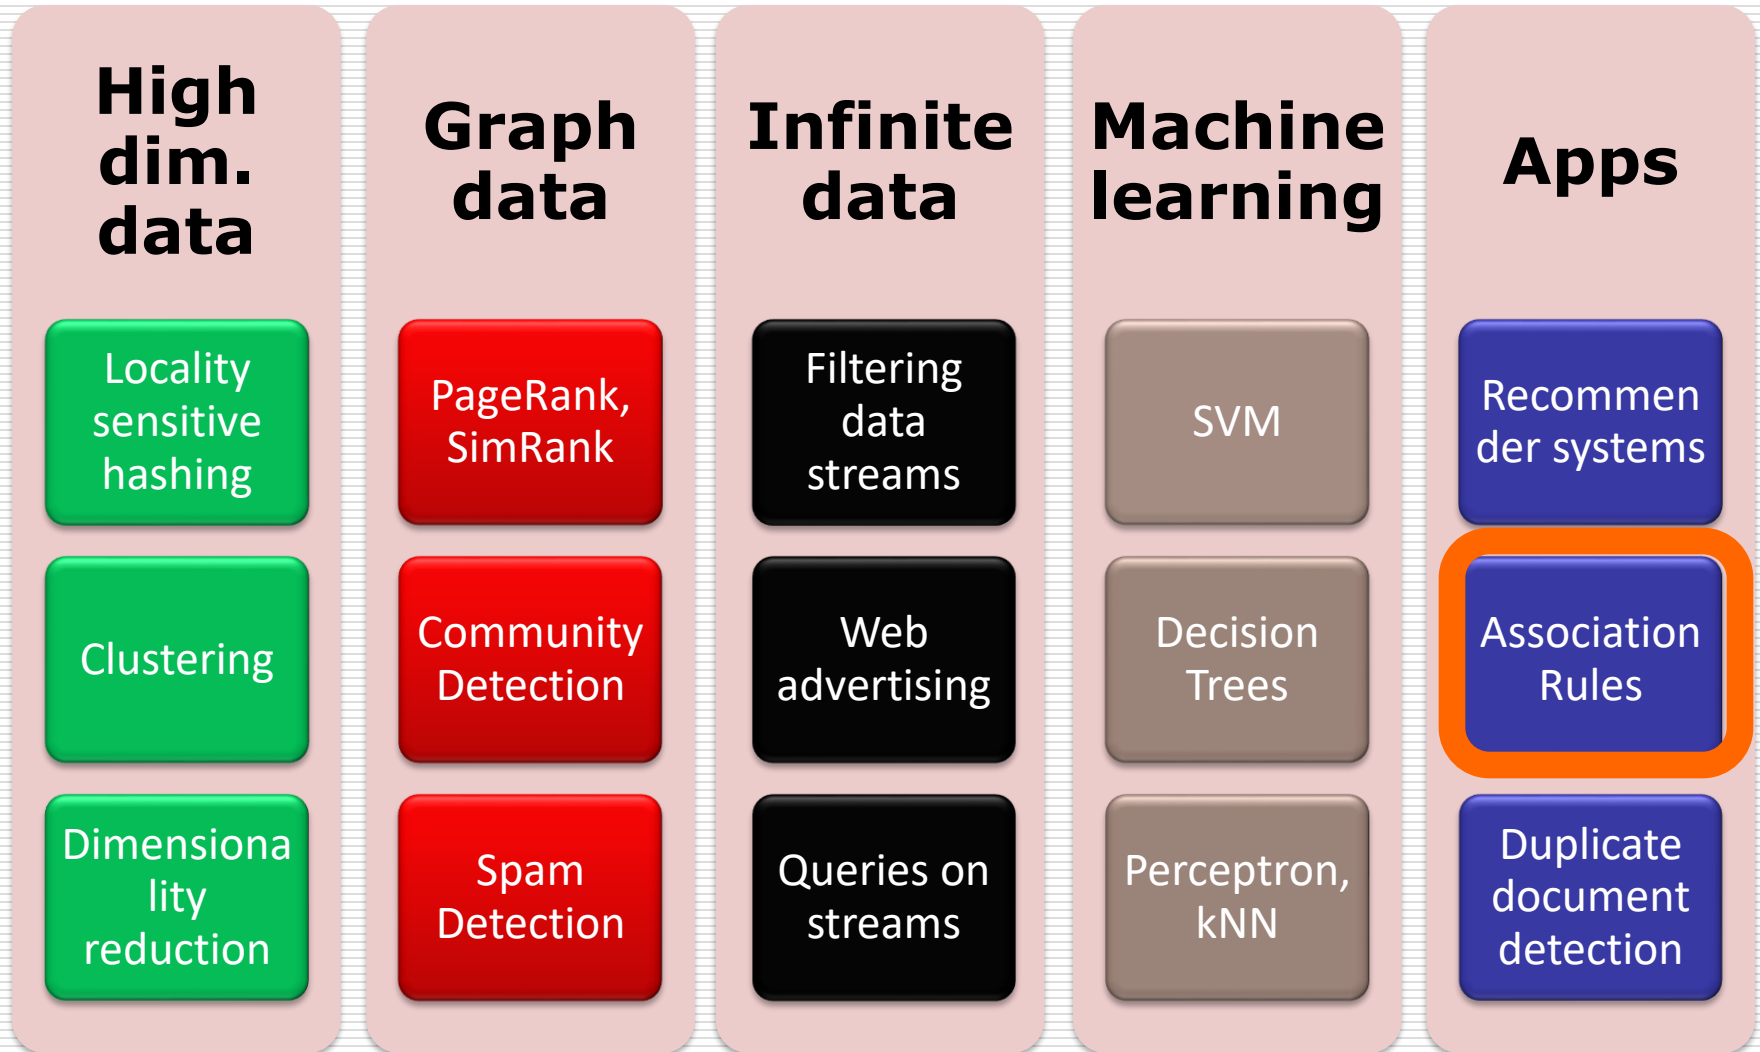


大数据计算及应用(二)

Association Rules and Frequent Pattern Mining

Slides adapted from <http://www.mmds.org>

Agenda



Association Rule

- Items frequently purchased together:

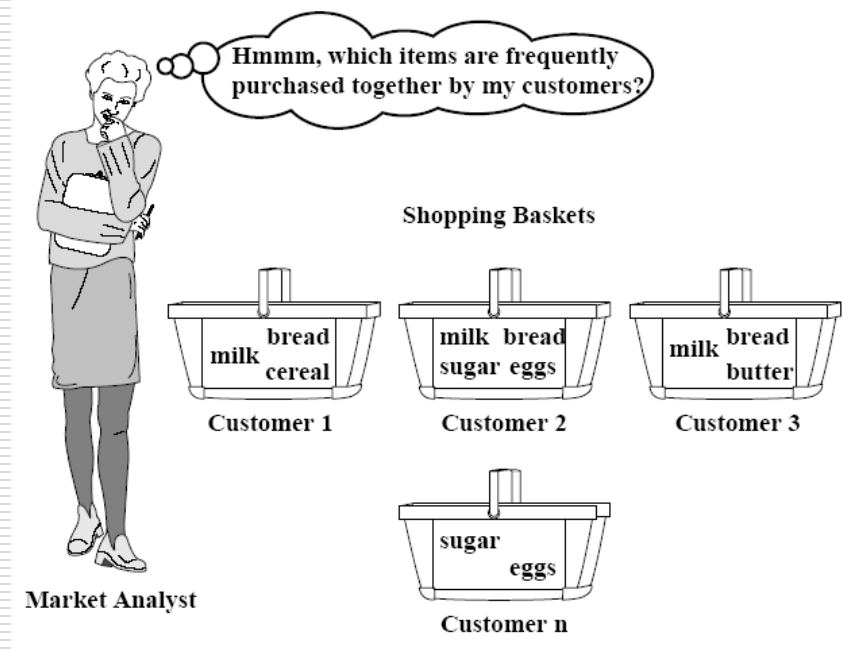
beer → diaper



- Uses:

- Placement
- Advertising
- Sales
- Coupons

- Objective: increase sales and reduce costs



The Market-Basket Model

- A large set of *items*, e.g., things sold in a supermarket
- A large set of *baskets*, each of which is a small set of the items, e.g., the things one customer buys on one day

<i>TID</i>	<i>Items</i>
1	Beer, Diaper, Milk
2	Coke, Diaper, Eggs
3	Beer, Coke, Diaper, Eggs
4	Coke, Eggs

The Market-Basket Model

- A general many-many mapping (association) between two kinds of things
 - But we ask about connections among “items” not “baskets”
- The technology focuses on **common events**, not rare events (“long tail”)

Applications – (1)

- **Items** = products; **baskets** = sets of products someone bought in one trip to the store
- **Example application**: given that many people buy beer and diapers together
 - Run a sale on diapers; raise price of beer
- Only useful if many buy diapers & beer

Applications – (2)

- **Items** = words; **Baskets** = Web pages;
- Unusual words appearing together in a large number of documents, e.g., “Brad” and “Angelina” may indicate an interesting relationship

Applications – (3)

- ❑ **Items** = sentences; **baskets** = documents containing those sentences
- ❑ Items that appear together too often could represent plagiarism

Association Rule Mining Applications

- ☐ Basket Data Analysis
- ☐ Genomic Data
- ☐ Telecommunication
- ☐ Credit Cards/ Banking Services
- ☐ Medical Treatments
- ☐ Web Personalization
- ☐ etc.

Scale of the Problem

- ❑ WalMart sells 100,000 items and can store billions of baskets
- ❑ The Web has billions of words and many billions of pages

Some Definition - Support

An itemset is **supported** by a basket (transaction) if it is included in the basket

Market-Basket transactions

<i>TID</i>	<i>Items</i>
1	Beer , Diaper , Milk
2	Coke, Diaper, Eggs
3	Beer , Coke, Diaper , Eggs
4	Coke, Eggs

<Beer, Diaper> is **supported** by basket 1, and 3, and its support is $2/4=50\%$.

Some Definition – Frequent Itemset

If the support of an itemset exceeds user specified *min_support* (threshold), this itemset is called a **frequent itemset (pattern)**.

Market-Basket transactions

<i>TID</i>	<i>Items</i>
1	Beer , Diaper , Milk
2	Coke, Diaper, Eggs
3	Beer , Coke, Diaper , Eggs
4	Coke, Eggs

min_support=50%

<Beer, Diaper> is a frequent itemset

<Beer, Milk> is not a frequent itemset

Association Rules

□ Association Rules:

If-then rules about the contents of baskets

□ $\{i_1, i_2, \dots, i_k\} \rightarrow j$ means: “if a basket contains all of i_1, \dots, i_k then it is *likely* to contain j ”

□ In practice there are many rules, want to find significant/interesting ones!

□ *Confidence* of this association rule is the probability of j given $I = \{i_1, \dots, i_k\}$

$$\text{conf}(I \rightarrow j) = \frac{\text{support}(I \cup j)}{\text{support}(I)}$$

Example: Confidence

$$T_1 = \{m, c, b\}$$

$$T_2 = \{m, p, j\}$$

$$T_3 = \{m, b\}$$

$$T_4 = \{c, j\}$$

$$T_5 = \{m, p, b\}$$

$$T_6 = \{m, c, b, j\}$$

$$T_7 = \{c, b, j\}$$

$$T_8 = \{b, c\}$$

□ Association rule: $\{m, b\} \rightarrow c$

■ $\text{Support}(m, b) = 4/8$, $\text{Support}(m, b, c) = 2/8$

■ **Confidence** = $2/4 = 0.5$


Association Rules Mining

- Question: "find all association rules with support $\geq s$ and confidence $\geq c$ "
- Hard part: finding the frequent itemsets

Computation Model

- Typically, data is kept in flat files rather than in a database system
 - Stored on disk
 - Stored basket-by-basket
 - Expand baskets into pairs, triples, etc. as you read baskets

File Organization



Example: items are positive integers, and boundaries between baskets are -1

Computation Model – (2)

- ❑ The true cost of mining disk-resident data is usually the **number of disk I/O's**
- ❑ In practice, association-rule algorithms read the data in *passes* – all baskets read in turn
- ❑ Thus, we measure the cost by the **number of passes** an algorithm takes

Main-Memory Bottleneck

- For many frequent-itemset algorithms, main memory is the critical resource
 - As we read baskets, we need to count something, e.g., occurrences of pairs
 - The number of different things we can count is limited by main memory
 - Swapping counts in/out is a disaster

Finding Frequent Pairs

- The hardest problem often turns out to be finding the frequent pairs
 - Why? Often frequent pairs are common, frequent triples are rare
 - Why? Probability of being frequent drops exponentially with size; number of sets grows more slowly with size
- We'll concentrate on pairs, then extend to larger sets

Naïve Algorithm

- ❑ Read file once, counting in main memory the occurrences of each pair
 - From each basket of n items, generate its $n(n-1)/2$ pairs by two nested loops
- ❑ Fails if $(\text{\#items})^2$ exceeds main memory
 - **Remember:** #items can be 100K (Wal-Mart) or 10B (Web pages)

Example: Counting Pairs

- ❑ Suppose 10^5 items
- ❑ Suppose counts are 4-byte integers
- ❑ Number of pairs of items: $10^5(10^5-1)/2 = 5*10^9$ (approximately)
- ❑ Therefore, $2*10^{10}$ (20 gigabytes) of main memory needed

Details of Main-Memory Counting

□ Two approaches:

(1) Count all pairs, using a triangular matrix

- requires only 4 bytes/pair


always assume integers are 4 bytes

(2) Keep a table of triples $[i, j, c]$ = "the count of the pair of items $\{i, j\}$ is c "


- requires 12 bytes

but only for those pairs with count > 0

Details of Main-Memory Counting



Method (1)



Method (2)

Comparing the Two Approaches

□ Approach 1: Triangular Matrix

- n = total number of items
- Count pair of items $\{i, j\}$ only if $i < j$
- Keep pair counts in lexicographic order:
 - $\{1,2\}, \{1,3\}, \dots, \{1,n\}, \{2,3\}, \{2,4\}, \dots, \{2,n\}, \{3,4\}, \dots, \{n-1,n\}$
- Pair $\{i, j\}$ is at position $(i-1)(n-i/2) + j - i$
- Total number of pairs $n(n-1)/2$; total bytes = $2n^2$
- **Triangular Matrix** requires 4 bytes per pair

□ Approach 2 uses **12 bytes** per occurring pair (*but only for pairs with count > 0*)

- Beats Approach 1 if less than **1/3** of possible pairs actually occur

Comparing the Two Approaches

□ Approach 1: Triangular Matrix

- n = total number items
- Count pair of items $\{i, j\}$ only if $i < j$

**Problem is if we
have too many
items so the pairs
do not fit into
memory.**

□ Approach 2:
(b)

- **Can we do better?**
possible pairs actually occur

...
 $s = 2n^2$

pair

Outline


- Association Rules
- Frequent Itemset Mining Algorithms
 - Apriori
 - FP-growth
- Sequential Pattern Mining Algorithms

Apriori Algorithm

□ Proposed by Rakesh Agrawal [VLDB'94]

□ Key idea:

- Candidate generation-and-test
- Anti-monotone property



<http://www.vldb.org> › conf PDF ⋮

Fast Algorithms for Mining Association Rules - VLDB ...

by R Agrawal · Cited by 28692 — We consider the problem of discovering **association rules** between items in a large database of sales transactions. We present two new **algorithms** for...
13 pages

Apriori Algorithm – (1)

- A two-pass approach called *Apriori* limits the need for main memory
- *Monotonicity* : if a set of items appears at least s times, so does every subset
 - *Contrapositive for pairs*: if item i does not appear in s baskets, then no pair including i can appear in s baskets

Apriori Algorithm – (2)

- **Pass 1:** Read baskets and count in main memory the occurrences of each item
 - Requires only memory proportional to #items
- Items that appear at least s times are the *frequent items*


Apriori Algorithm – (3)

- **Pass 2:** Read baskets again and count in main memory only those pairs both of which were found in Pass 1 to be frequent
- Requires memory proportional to square of *frequent* items only (for counts), plus a list of the frequent items (so you know what must be counted)

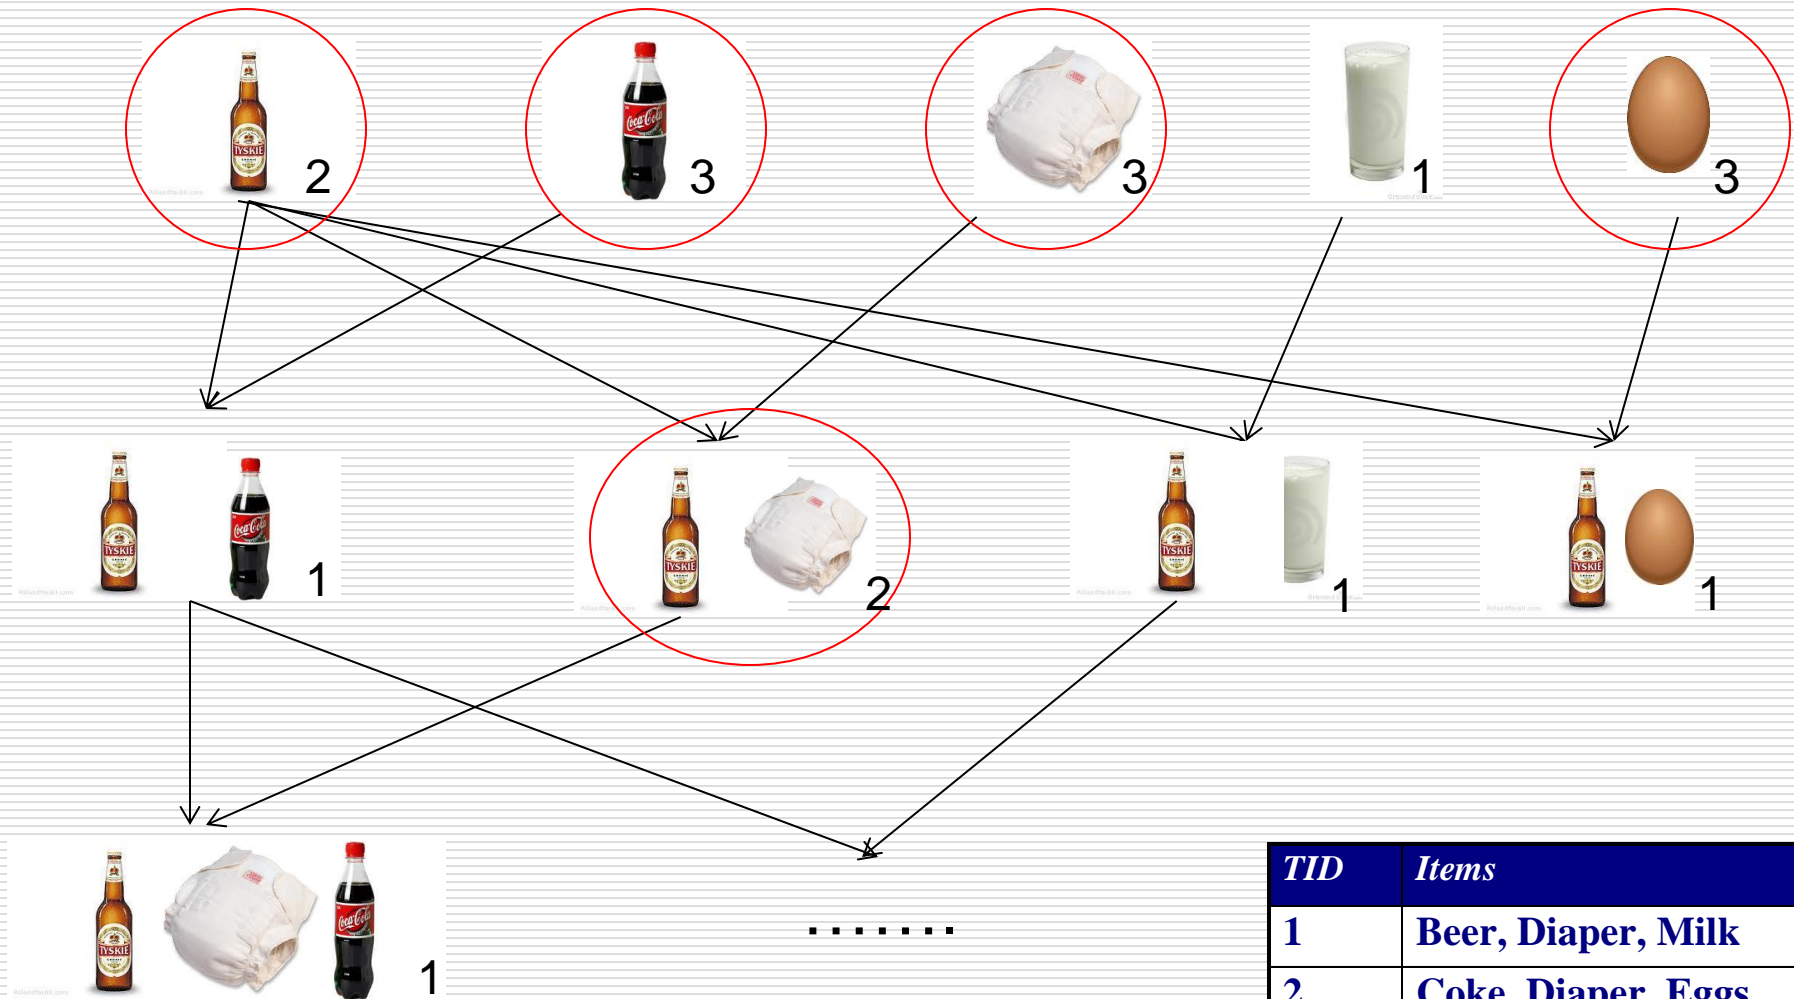
Apriori Algorithm

Market-Basket transactions

<i>TID</i>	<i>Items</i>
1	Beer, Diaper, Milk
2	Coke, Diaper, Eggs
3	Beer, Coke, Diaper, Eggs
4	Coke, Eggs



Naive Algorithm




Sup_{min}=2

TID	Items
1	Beer, Diaper, Milk
2	Coke, Diaper, Eggs
3	Beer, Coke, Diaper, Eggs
4	Coke, Eggs

Apriori Algorithm

- Anti-monotone property: If an itemset is not frequent, then any of its superset is not frequent



Sup_{min}=2

<i>TID</i>	<i>Items</i>
1	Beer, Diaper, Milk
2	Coke, Diaper, Eggs
3	Beer, Coke, Diaper, Eggs
4	Coke, Eggs

Apriori Algorithm

$\text{Sup}_{\min} = 2$

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E


1st scan

C_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

L_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3



C_2

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

2nd scan

C_2

Itemset
{A, B}
{A, C}
{A, E}
{B, C}
{B, E}
{C, E}

L_2

Itemset	sup
{A, C}	2
{B, C}	2
{B, E}	3
{C, E}	2



C_3

Itemset
{B, C, E}

3rd scan


L_3

Itemset	sup
{B, C, E}	2

Apriori Algorithm

1. C_1 = Itemsets of size one in I ;
2. Determine all large itemsets of size 1, L_1 ;
3. $i = 1$;
4. Repeat
5. $i = i + 1$;
6. C_i = Apriori-Gen(L_{i-1});
7. Count C_i to determine L_i ;
8. until no more large itemsets found;

Frequent Itemset Property



Drawbacks of Apriori

- Multiple scans of transaction database
 - Multiple database scans are **costly**
- Huge number of candidates
 - To find frequent itemset $i_1 i_2 \dots i_{100}$
 - # of scans: **100**
 - # of Candidates: $2^{100}-1 = 1.27*10^{30}$

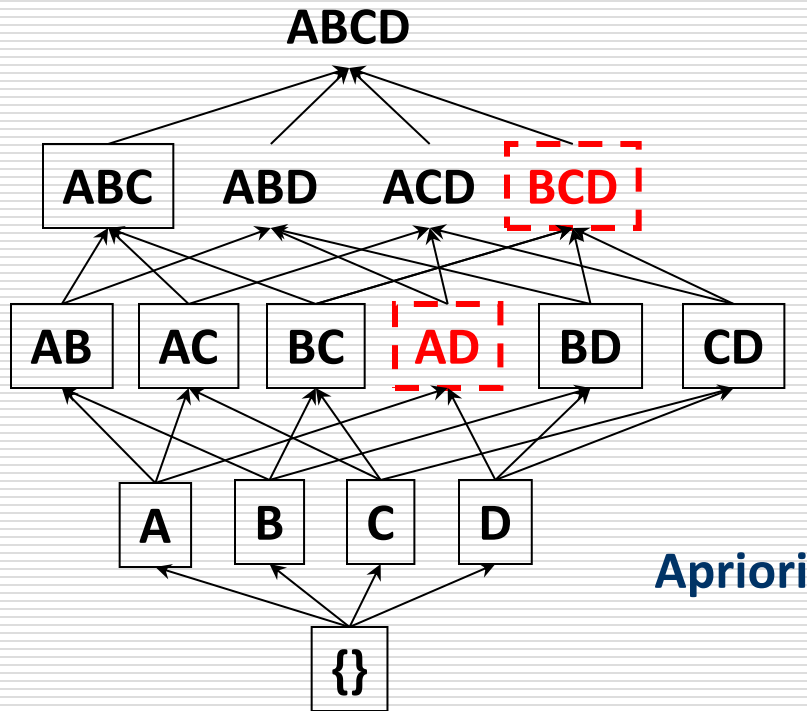
Improving Apriori: General Ideas

- ❑ Reduce passes of transaction database scans
- ❑ Shrink number of candidates
- ❑ Facilitate support counting of candidates

Improving Apriori's Efficiency

- ❑ **Hash-based itemset counting:** A k -itemset whose corresponding hashing bucket count is below the threshold cannot be frequent
- ❑ **Transaction reduction:** A transaction that does not contain any frequent k -itemset is useless in subsequent scans
- ❑ **Partitioning:** Any itemset that is potentially frequent in DB must be frequent in at least one of the partitions of DB
- ❑ **Sampling:** mining on a subset of given data, need a lower support threshold + a method to determine the completeness
- ❑ **Dynamic itemset counting:** add new candidate itemsets immediately (unlike Apriori) when all of their subsets are estimated to be frequent



DIC: Reduce Number of Scans



Itemset lattice

S. Brin, R. Motwani, J. Ullman, and S. Tsur. Dynamic itemset counting and implication rules for market basket data. In *SIGMOD'97*

- Once both A and D are determined frequent, the counting of AD begins
- Once all length-2 subsets of BCD are determined frequent, the counting of BCD begins



FP-growth Algorithm

FP-Growth

- Proposed by Jiawei Han [*SIGMOD'00*]
- Uses the Apriori pruning principle
- Scan DB only twice
 - Once to find frequent 1-itemset (single item pattern)
 - Once to construct FP-tree (prefix tree, Trie), the data structure of FP-growth



Mining Frequent Patterns without Candidate Generation

- Compress a large database into a compact, Frequent-Pattern tree (FP-tree) structure
 - highly condensed, but complete for frequent pattern mining
 - avoid costly database scans
- Develop an efficient, FP-tree-based frequent pattern mining method
 - A divide-and-conquer methodology: decompose mining tasks into smaller ones
 - Avoid candidate generation: sub-database test only!

FP-Growth

$\text{Sup}_{\min} = 2$

<i>TID</i>	<i>Items bought</i>
10	{f, a, c, d, g, i, m, p}
20	{a, b, c, f, l, m, o}
30	{b, f, h, j, o, w}
40	{b, c, k, s, p}
50	{a, f, c, e, l, p, m, n}


Header Table

Item frequency

<i>f</i>	4
<i>c</i>	4
<i>a</i>	3
<i>b</i>	3
<i>m</i>	3
<i>p</i>	3



<i>TID</i>	<i>(ordered) frequent items</i>
10	{f, c, a, m, p}
20	{f, c, a, b, m}
30	{f, b}
40	{c, b, p}
50	{f, c, a, m, p}



FP-Growth

$\text{Sup}_{\min} = 2$


<i>TID</i>	<i>Items bought</i>
10	{f, a, c, d, g, i, m, p}
20	{a, b, c, f, l, m, o}
30	{b, f, h, j, o, w}
40	{b, c, k, s, p}
50	{a, f, c, e, l, p, m, n}

Header Table

<i>Item</i>	<i>frequency</i>
<i>f</i>	4
<i>c</i>	4
<i>a</i>	3
<i>b</i>	3
<i>m</i>	3
<i>p</i>	3



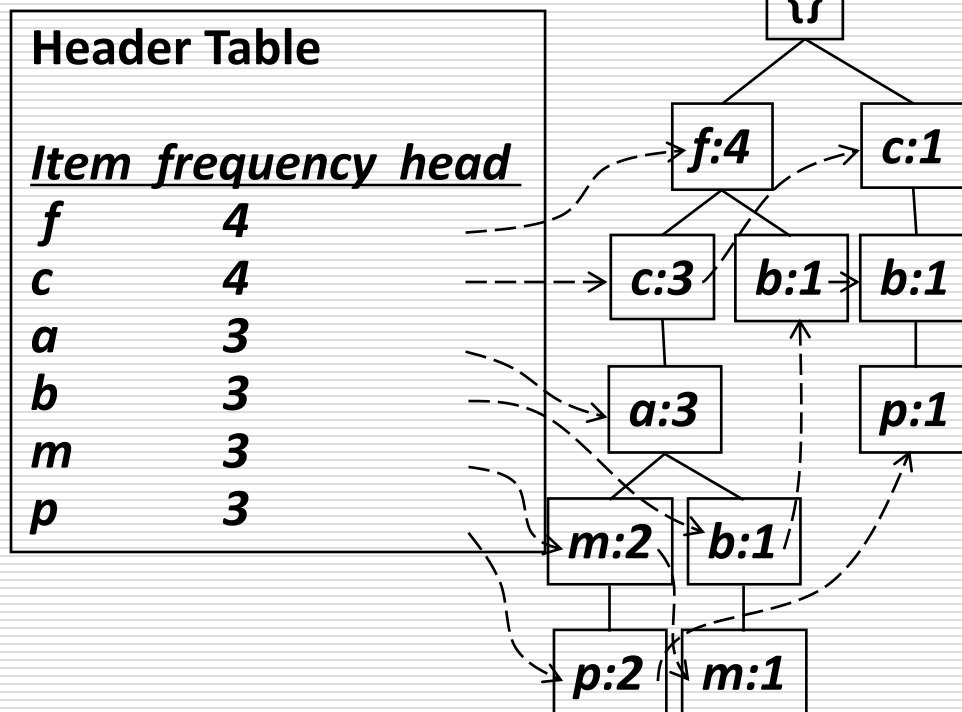
<i>TID</i>	<i>(ordered) frequent items</i>
10	{f, c, a, m, p}
20	{f, c, a, b, m}
30	{f, b}
40	{c, b, p}
50	{f, c, a, m, p}



FP-Growth

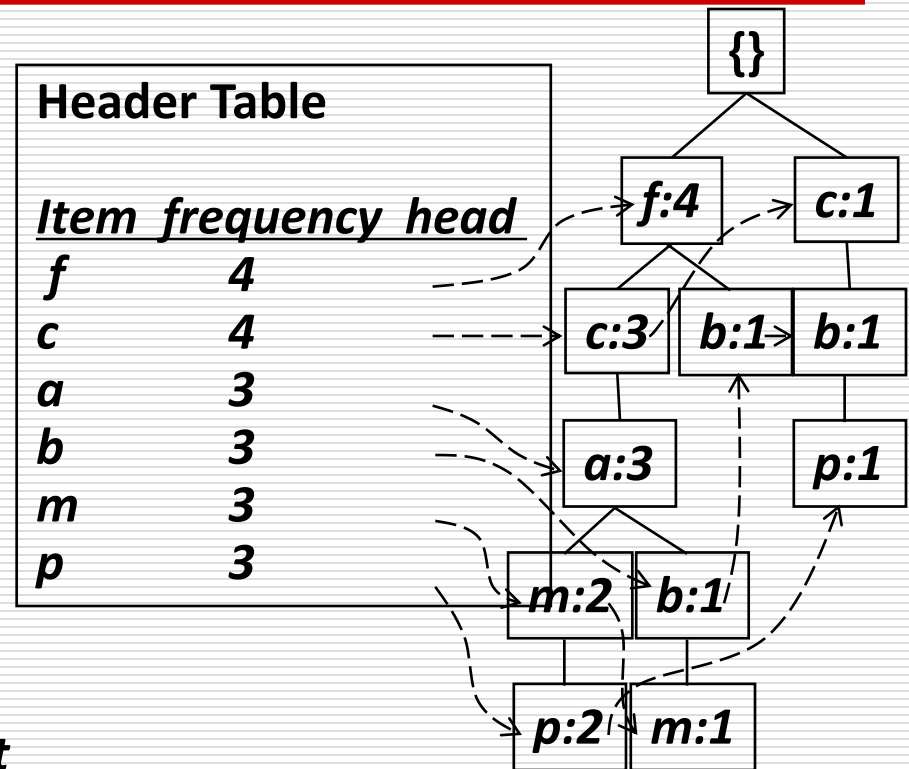
$$\text{Sup}_{\min} = 2$$

<i>TID</i>	<i>(ordered) frequent items</i>
10	{ <i>f</i> , <i>c</i> , <i>a</i> , <i>m</i> , <i>p</i> }
20	{ <i>f</i> , <i>c</i> , <i>a</i> , <i>b</i> , <i>m</i> }
30	{ <i>f</i> , <i>b</i> }
40	{ <i>c</i> , <i>b</i> , <i>p</i> }
50	{ <i>f</i> , <i>c</i> , <i>a</i> , <i>m</i> , <i>p</i> }



FP-Growth

$$\text{Sup}_{\min} = 2$$



Conditional pattern bases

Item cond. pattern base freq. itemset

<i>p</i>	<i>fcam:2, cb:1</i>	<i>fp, cp, ap, mp, fcp, fap, fmp, cap, cmp, amp, camp, facp, fcmp, famp, fcamp</i>
<i>m</i>	<i>fca:2, fcab:1</i>	<i>fm, cm, am, fcm, fam, cam, fcam</i>
<i>b</i>	<i>fca:1, f:1, c:1</i>	...
<i>a</i>	<i>fc:3</i>	...
<i>c</i>	<i>f:3</i>	...

Why Is Frequent Pattern Growth Fast?

- The performance study shows
 - FP-growth is faster than Apriori (in most cases), and is also faster than tree-projection (an order of magnitude on some datasets)
- Reasoning
 - No candidate generation (claimed by the authors)
 - Use compact data structure
 - Eliminate repeated database scan
 - Basic operation is counting and FP-tree building

Extension of Association Rule Mining

- ❑ Association rule mining has been extensively studied in the data mining community.
- ❑ There are many efficient algorithms and model variations.
- ❑ Other related work includes
 - Multi-level or generalized rule mining
 - Sequential pattern mining
 - Constrained rule mining
 - Incremental rule mining
 - Maximal and closed frequent itemset mining
 - Numeric association rule mining
 - Rule interestingness and visualization
 - Parallel algorithms
 - ...

Extension of Association Rule Mining

- ❑ Association rule mining has been extensively studied in the data mining community.
- ❑ There are many efficient algorithms and model variations.
- ❑ Other related work includes
 - Multi-level or generalized rule mining
 - [Sequential pattern mining](#)
 - Constrained rule mining
 - Incremental rule mining
 - Maximal and closed frequent itemset mining
 - Numeric association rule mining
 - Rule interestingness and visualization
 - Parallel algorithms
 - ...

Applications

Web Images Maps News Shopping Gmail more ▾

Google™ **sequence mining** Search Advanced Search

Web

Sequence mining - Wikipedia, the free encyclopedia
Sequence mining is concerned with finding statistically relevant patterns between data examples where the values are delivered in a sequence. ...
en.wikipedia.org/wiki/Sequence_mining - 15k - [Cached](#) - [Similar pages](#)

Publications by Mohammed Javeed Zaki
Mohammed J. Zaki, **Sequence Mining** in Categorical Domains: Algorithms and ...
Mohammed J. Zaki, **Parallel Sequence Mining** on Shared-Memory Machines, ...
www.cs.rpi.edu/~zaki/papers.html - 66k - [Cached](#) - [Similar pages](#)

[PDF] 7. Sequence Mining ← Clicked page
File Format: Microsoft Powerpoint - [View as HTML](#)
Sequence Mining. Sequences and Strings. Recognition with Strings. MM & HMM.
Sequence Association Rules. 7/03. **Data Mining** - Sequences ...
www.cs.wright.edu/~gdong/mining03/course_slides/C7Seq.ppt - [Similar pages](#)

Incremental and Interactive Sequence Mining - Parthasarathy ...
Interactive **Sequence Discovery** by Incremental **Mining** - Ming-Yen Lin And (Correct) ... 0.3:
Parallel Sequence Mining on Shared-Memory Machines - Zaki (2000) ...
citeseer.ist.psu.edu/235156.html - 25k - [Cached](#) - [Similar pages](#)

Sequence Mining in Categorical Domains: Incorporating Constraints ...
We present cSPADE, an efficient algorithm for **mining** frequent sequences considering a variety of syntactic constraints.
citeseer.ist.psu.edu/616916.html - 24k - [Cached](#) - [Similar pages](#)
[More results from citeseer.ist.psu.edu »](#)

Web Images Maps News Shopping Gmail more ▾

Google™ **sequential pattern mining** Search Advanced Search

Web

SPAM: Sequential Pattern Mining ← Re-query
SPAM: **Sequential Pattern Mining**. SPAM is a new algorithm for finding all frequent sequences within a transactional database. The algorithm is especially ...
himalaya-tools.sourceforge.net/Spam/ - 7k - [Cached](#) - [Similar pages](#)

[PDF] Multi-dimensional Sequential Pattern Mining
File Format: PDF/Adobe Acrobat - [View as HTML](#)
Multi-dimensional **Sequential Pattern Mining**. Helen Pinto. Jiawei Han. Jian Pei
SEQUENTIAL PATTERN MINING. "m1Z@h"kn"r4j&o1l(mb f({Hm1r,h"ka(m1AkEnAo ...
www.sal.cs.uiuc.edu/~hanj/pdf/mdseq01.pdf - [Similar pages](#)

[PDF] Sequential Pattern Mining Algorithms: Trade-offs between Speed and ...
File Format: PDF/Adobe Acrobat - [View as HTML](#)
In this paper, we study the problem of **sequential pattern mining**. ... goal of **sequential pattern mining** is to discover all frequent sequences of itemsets in ...
hms.liacs.nl/mgts2004/papers/antunes.pdf - [Similar pages](#)


[PDF] Sequential Pattern Mining: A Survey
File Format: PDF/Adobe Acrobat - [View as HTML](#)
those techniques web **mining** and **sequential pattern mining** are also well
Sequential patterns: sequential pattern mining is trying to find the ...
www.icar.cnr.it/.../2006/datamining/Esami2006/ArticoliSelezionatiDM/SEMINARI/Mining%20Data%20Streams/SQM.pdf - [Similar pages](#)

Mining Sequential Patterns by Pattern-Growth: The PrefixSpan ...
Sequential pattern mining is an important data **mining** problem with broad applications. However, it is also a difficult problem since the **mining** may have to ...
citeseer.ist.psu.edu/655475.html - 25k - [Cached](#) - [Similar pages](#)

Mining Sequential Patterns - Agrawal, Srikant (ResearchIndex)
We introduce the problem of **mining sequential patterns** over such databases. We present three algorithms to solve this problem, and empirically evaluate ...
citeseer.ist.psu.edu/agrawal95mining.html - 25k - [Cached](#) - [Similar pages](#)

Examples of Sequence Data

Sequence Database	Sequence	Element (Transaction)	Event (Item)
Customer	Purchase history of a given customer	A set of items bought by a customer at time t	Books, diary products, CDs, etc
Web Data	Browsing activity of a particular Web visitor	A collection of files viewed by a Web visitor after a single mouse click	Home page, index page, contact info, etc
Event data	History of events generated by a given sensor	Events triggered by a sensor at time t	Types of alarms generated by sensors
Genome sequences	DNA sequence of a particular species	An element of the DNA sequence	Bases A,T,G,C



Formal Definition of a Sequence

- A sequence is an ordered list of elements (transactions)

$$S = \langle e_1 e_2 e_3 \dots \rangle$$

- Each element contains a collection of items

$$e_i = \{i_1, i_2, \dots, i_k\}$$

- Each element is attributed to a specific time

- A k-sequence is a sequence that contains k items

Formal Definition of a Subsequence

- A sequence $\langle a_1 a_2 \dots a_n \rangle$ is contained in another sequence $\langle b_1 b_2 \dots b_m \rangle$ ($m \geq n$) if there exist integers $i_1 < i_2 < \dots < i_n$ such that $a_1 \subseteq b_{i_1}, a_2 \subseteq b_{i_2}, \dots, a_n \subseteq b_{i_n}$

Data sequence	Subsequence	Contain?
$\langle \{2,4\} \{3,5,6\} \{8\} \rangle$	$\langle \{2\} \{3,5\} \rangle$	Yes
$\langle \{1,2\} \{3,4\} \rangle$	$\langle \{1\} \{2\} \rangle$	No
$\langle \{2,4\} \{2,4\} \{2,5\} \rangle$	$\langle \{2\} \{4\} \rangle$	Yes

- The support of a subsequence w is defined as the fraction of data sequences that contain w
- A *sequential pattern* is a frequent subsequence (i.e., a subsequence whose support is $\geq \text{minsup}$)

Sequential Pattern Mining: Definition

□ Given:


- a database of sequences
- a user-specified minimum support threshold, *minsup*

□ Task:

- Find all subsequences with support $\geq \textit{minsup}$

Sequential Pattern Mining: Challenge

- How many k-subsequences can be extracted from a given n-sequence?



Answer:

$$\binom{n}{k} = \binom{9}{4} = 126$$

Outline

- Association Rules
- Frequent Itemset Mining Algorithms
- Sequential Pattern Mining Algorithms
 - GSP
 - SPADE
 - SPAM

GSP (Generalized Sequential Pattern Mining)

- Proposed by Srikant and Agrawal [*EDBT'96*]
- Uses the Apriori pruning principle

Finding Length-1 Sequential Patterns

□ Initial candidates:

■ $\langle a \rangle, \langle b \rangle, \langle c \rangle, \langle d \rangle, \langle e \rangle, \langle f \rangle, \langle g \rangle, \langle h \rangle$

□ Scan database once, count support for candidates

$min_sup = 2$

Seq. ID	Sequence
10	$\langle (bd)cb(ac) \rangle$
20	$\langle (bf)(ce)b(fg) \rangle$
30	$\langle (ah)(bf)abf \rangle$
40	$\langle (be)(ce)d \rangle$
50	$\langle a(bd)bcb(ade) \rangle$

Cand	Sup
$\langle a \rangle$	3
$\langle b \rangle$	5
$\langle c \rangle$	4
$\langle d \rangle$	3
$\langle e \rangle$	3
$\langle f \rangle$	2
$\langle g \rangle$	1
$\langle h \rangle$	1

Generating Length-2 Candidates

51 length-2
Candidates

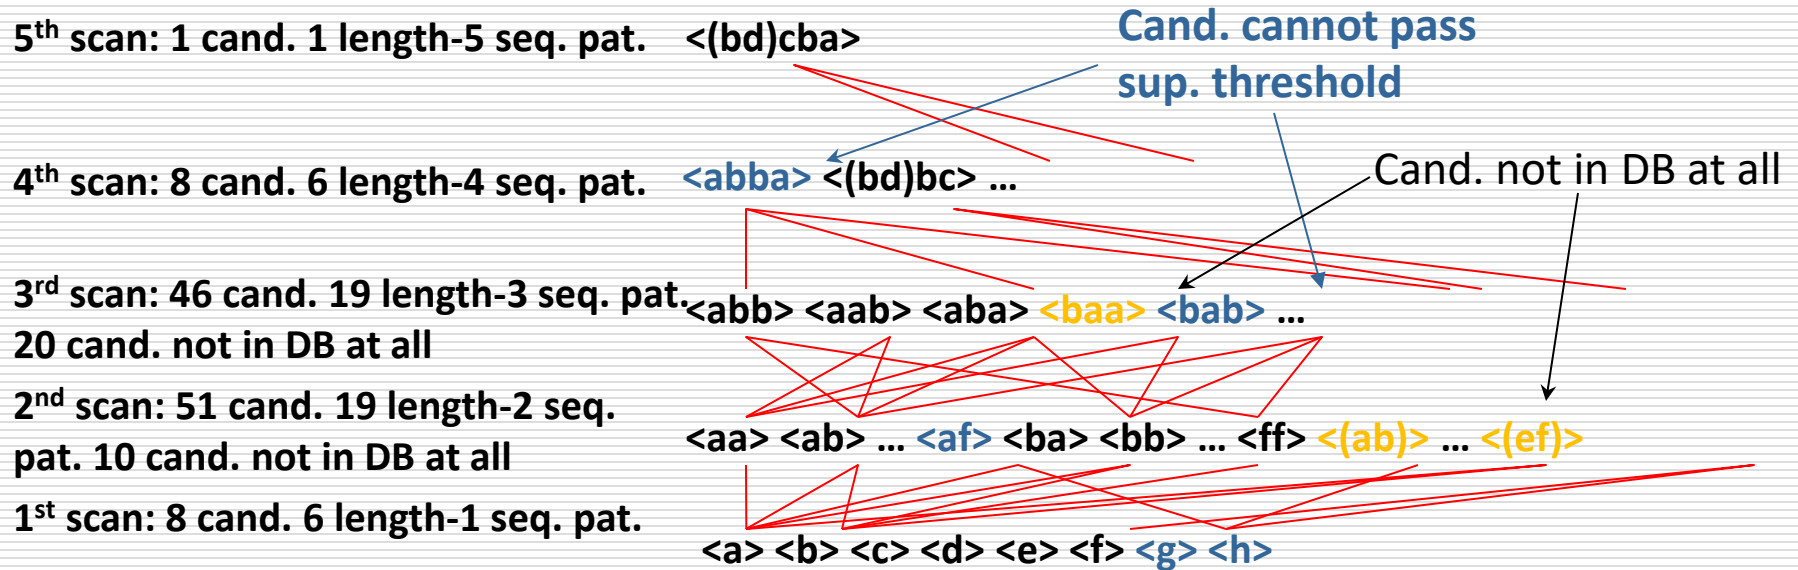
	<a>		<c>	<d>	<e>	<f>
<a>	<aa>	<ab>	<ac>	<ad>	<ae>	<af>
	<ba>	<bb>	<bc>	<bd>	<be>	<bf>
<c>	<ca>	<cb>	<cc>	<cd>	<ce>	<cf>
<d>	<da>	<db>	<dc>	<dd>	<de>	<df>
<e>	<ea>	<eb>	<ec>	<ed>	<ee>	<ef>
<f>	<fa>	<fb>	<fc>	<fd>	<fe>	<ff>

	<a>		<c>	<d>	<e>	<f>
<a>		<(ab)>	<(ac)>	<(ad)>	<(ae)>	<(af)>
			<(bc)>	<(bd)>	<(be)>	<(bf)>
<c>				<(cd)>	<(ce)>	<(cf)>
<d>					<(de)>	<(df)>
<e>						<(ef)>
<f>						

Without Apriori
property,
 $8*8+8*7/2=92$
candidates

Apriori prunes
44.57%
candidates

GSP Mining Process



min_sup = 2

Seq. ID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

GSP Algorithm

- Take sequences in form of $\langle x \rangle$ as length-1 candidates
- Scan database once, find F_1 , the set of length-1 sequential patterns
- Let $k=1$; while F_k is not empty do
 - Form C_{k+1} , the set of length- $(k+1)$ candidates from F_k ;
 - If C_{k+1} is not empty, scan database once, find F_{k+1} , the set of length- $(k+1)$ sequential patterns
 - Let $k=k+1$;

GSP Algorithm

- Benefits from the Apriori pruning
 - Reduces search space
- Bottlenecks
 - Scans the database multiple times
 - Generates a huge set of candidate sequences

SPADE Algorithm

SPADE Algorithm

- ❑ Proposed by Zaki *et al.* [*MLJ'01*]
- ❑ Candidate generation-and-test
- ❑ Vertical ID-list data representation based on Lattice-theory
- ❑ Counting support through temporal joins
- ❑ Reduced I/O costs (three DB scans)

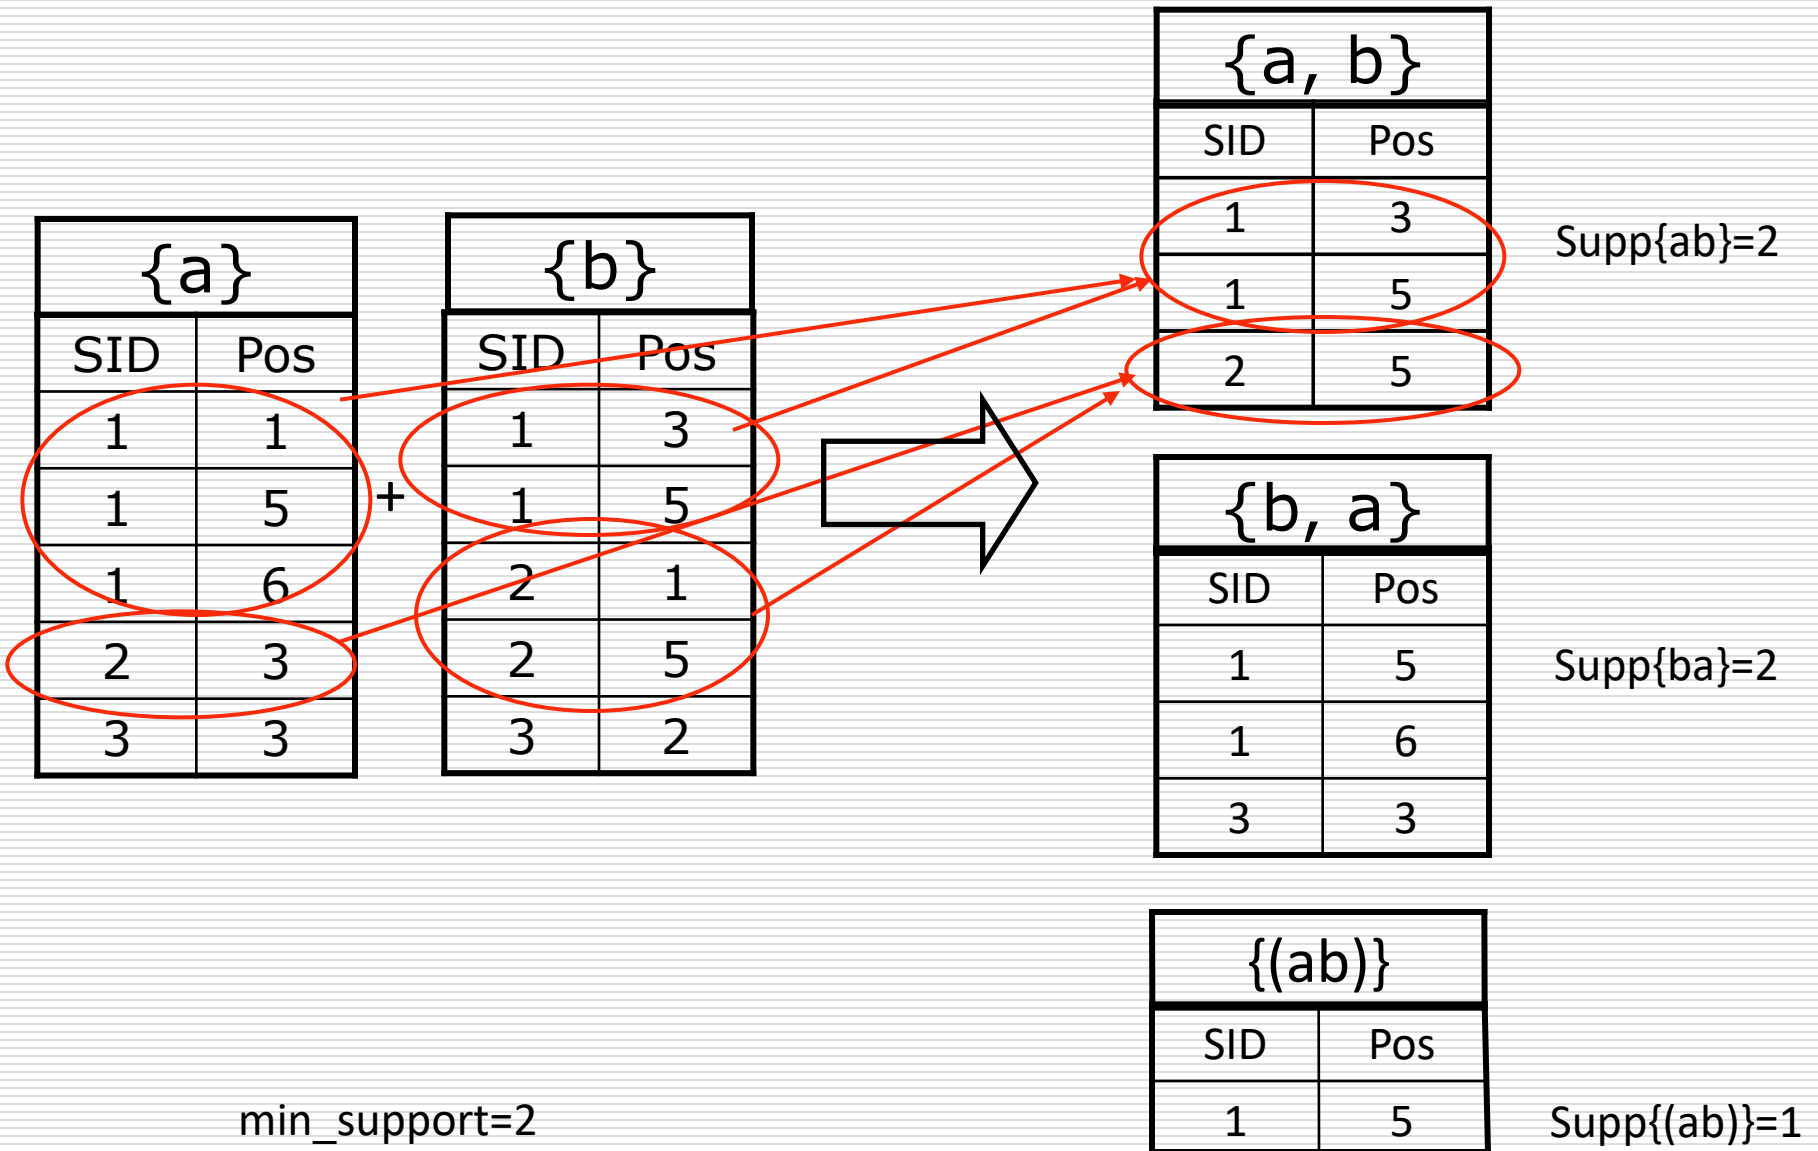
SPADE Algorithm

ID	Data Sequence
1	<a c (b c) d (a b c) a d >
2	<b (c d) a c (b d) >
3	<d (b c) (a c) (c d) >

ID-List DB

a		b		c		d	
SID	Pos	SID	Pos	SID	Pos	SID	Pos
1	1	1	3	1	2	1	4
1	5	1	5	1	3	1	7
1	6	2	1	1	5	2	2
2	3	2	5	2	2	2	5
3	3	3	2	2	4	3	1
				3	2	3	4
				3	3		
				3	4		

Temporal Joins



SPAM Algorithm

SPAM Algorithm

- ❑ Proposed by Ayres *et al.* [*KDD'02*]
- ❑ Key idea based on SPADE
- ❑ Using bitmap data representation
- ❑ Faster than SPADE yet space consuming

SPAM Algorithm

ID	Data Sequence
10	<a c (b c) d (a b c) a d>
20	<b (c d) a c (b d)>
30	<d (b c) (a c) (c d)>

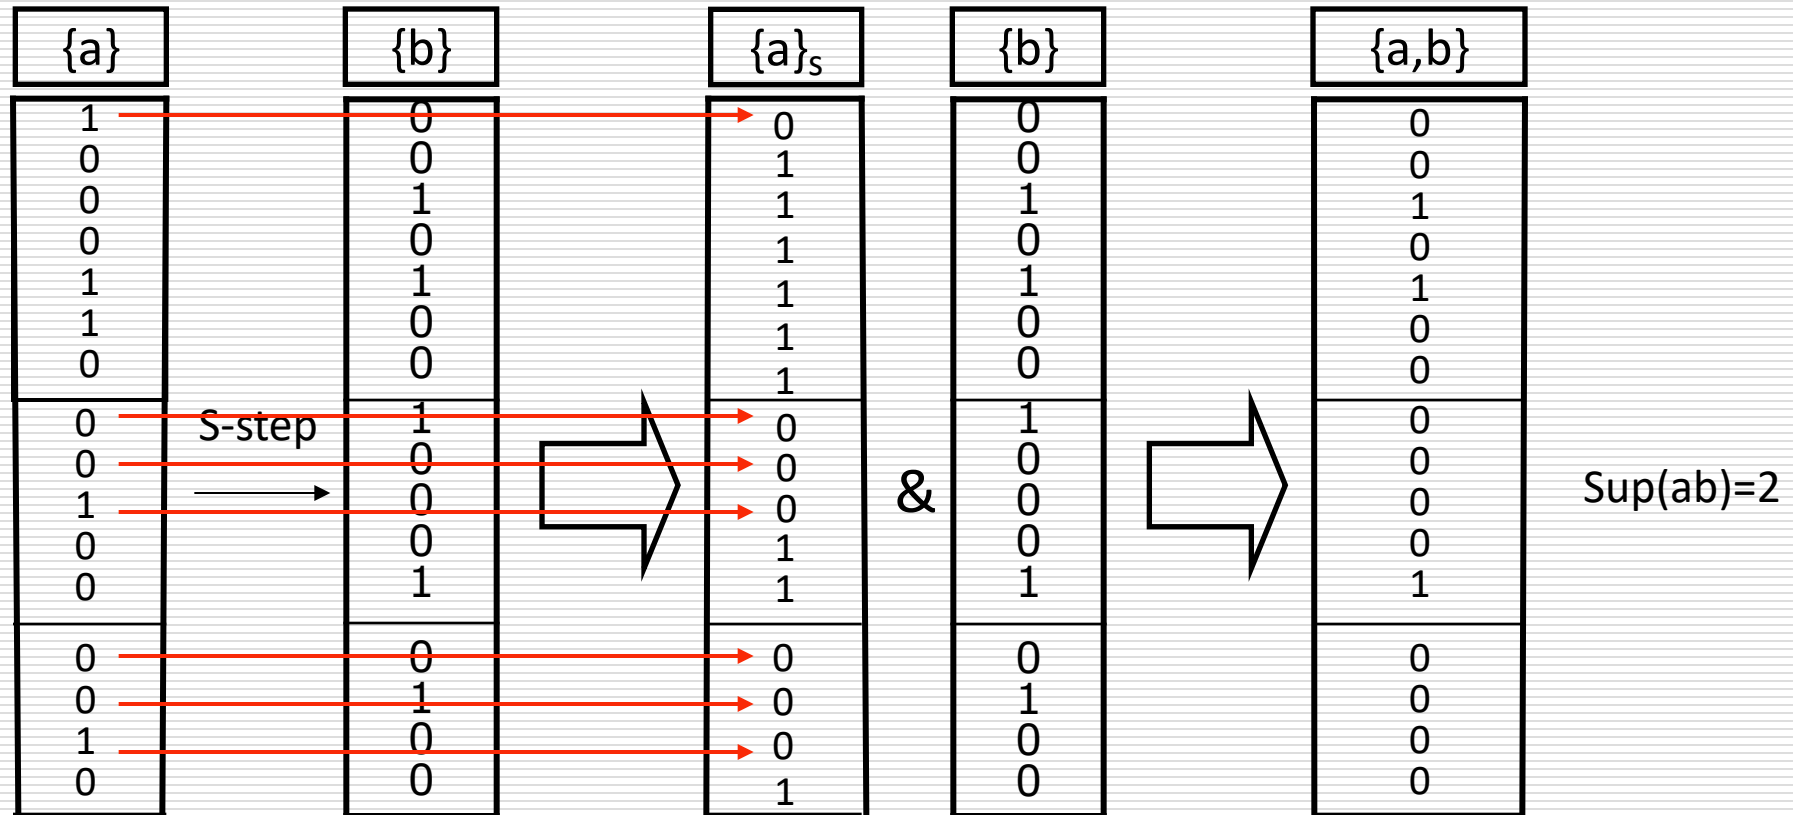
SID EID

10 1
10 2
10 3
10 4
10 5
10 6
10 7
20 1
20 2
20 3
20 4
20 5
30 1
30 2
30 3
30 4

{a}	{b}	{c}	{d}
1	0	0	0
0	0	1	0
0	1	1	0
0	0	0	1
1	1	1	0
1	0	0	0
0	0	0	1
0	1	0	0
0	0	1	1
0	0	0	0
0	1	0	1
0	0	0	1
0	1	1	0
1	0	1	0
0	0	1	1

SPAM Temporal Joins

Sequence extended step:



min_support=2

Problem of SPAM

- Bitmap representation is space consuming
i.e., data is commonly very sparse

Acknowledgement

- Slides are adapted from:
 - Prof. Jeffrey D. Ullman
 - Dr. Jure Leskovec
 - Dr. Wujun Li